Homework 2

25: not

Theory

Book problems 25, 29, 30, 31, 32

Programming

The original idea for this assignment came from Nir Shavit. Lamont Samuels developed helper code for C for the UChicago version of this assignment.

The Programming Environment

We will continue programming in C with POSIX threads (pthreads). **Your final codes for this assignment should not make use of pthreads synchronization routines**. For this assignment, synchronization will be done through a Lamport Queue. You will need an atomic builtin from gcc (a memory barrier or memory fence) and the use of the volatile keyword to properly implement this data structure, however.

Programming Projects Background

The remaining programming assignments are based on building a firewall: we will see a stream of *packets* from various source addresses (represented as integers) and perform several operations before sending them to the destination address (also integers) specified in the header of the packet. By default, packets will be *data* packets, messages from source address to destination address with some variable-length payload. Our firewall will perform two primary functions: 1) enforce access controls and 2) filter the data payloads for evidence of *bad guys* (that the packet is from a suspicious source).

The Programming Assignment

This assignment is all about work distribution and load balancing. Suppose that we have $n-1$ different sources, each of which will supply T packets (n and T are input parameters). Your goal is to calculate a checksum on all packets. You'll be provided with a data object and associated functions that calculate the checksum and you will also be provided with an object which creates the packets (the *Packet Generator*). In order to get a parallel speedup, you will have a single thread (the *Dispatcher*) retrieve the packets from the packet generator, but distribute the work of calculating the checksums across $n-1$ threads (the *Workers*). Because we tie the number of threads to the number of sources, we can use a Lamport queue (Figure 3.3 "WaitFreeQueue" in the text—also called a single-reader / single-writer queue) of depth D (an input parameter) to buffer packets for each thread. That is, each worker has a dedicated Lamport queue which is written only by the Dispatcher and read only by the associated Worker, thus requiring no locks. Provided the typical work to calculate a checksum is large relative to the work to retrieve the packet and write it into the Lamport queue, we should see a reasonable opportunity for parallel speedup. We will build three different versions of the code:

- SERIAL The application, as described above, where a thread loops through each source and each packet to calculate the checksums serially.

- PARALLEL The application, as described above, implemented with $n-1$ workers calculating the checksums in parallel.

- SERIAL-QUEUE The application, as described above, in which a thread proceeds in much the same way as in SERIAL, but each packet from the ith source must be written to and then subsequently read from the ith Lamport queue. In particular, all data structures used in the PARALLEL version should be instantiated in this version, the only difference being that all code is performed by a single thread. This configuration allows us to ascertain (almost all—what is missing?) the overhead in additional work we incur by parallelizing the application and communicating through the queue.

Load Balancing

Assuming that the work is dominated by the checksum calculation, we might indeed get decent parallel speedup. However, this speedup also relies on how the distribution of work falls to each worker, since we have a fixed assignment of worker to packet source. That is, if one worker has considerably more work than the average worker, the overall runtime will suffer. To explore this effect, we will use three different sources of packets: *Constant*, *Uniform*, and *Exponential.* Constant packets have the same amount of work per packet, Uniform packets have the property that the expected work for every packet from every source is identical and distributed according to a uniform random variable (uniform over $[0, 2W]$), where the expected work, W, is an input parameter. Exponential packets, by contrast, have for each packet coming from source i an expected work equal to μ_i, which is distributed according to an exponential distribution with parameter $\lambda = \frac{1}{W}$. Furthermore, the packets from source i have work that is distributed according to an exponential distribution with parameter $\lambda_i = \frac{1}{\mu_i}$. The difference in average work per source and the added variability of the exponential distribution (vs. the uniform distribution) both give rise to an imbalance of work across the $n - 1$ workers.

Code

The math here???

You should create a directory in svn and expand the tarball we provided on the class website.

In particular, the following data types are provided:

- PACKETSOURCE_T is a random packet generator (see packetsource.h) which is instantiated with three parameters:

    - *mean* is the expected amount of work that will need to be performed per packet.
    - *numSources* is the number of packet sources.
    - *seed* initializes the internal random number generator - it can be set to the trial number if you wish to corroborate your packet checksums across runs of SERIAL, SERIAL-QUEUE and PARALLEL.

    In addition to creation and deletion, it exports three functions, GETCONSTANTPACKET, GETUNIFORMPACKET, and GETEXPONENTIALPACKET, which, given an input source number, returns the next packet from that source. *Note: the uniform and exponential packet streams are independent for each source - so the sequence of packets from each source is deterministic, regardless of the order that they are taken among the sources.*

- STOPWATCH is a simple timer (see stopwatch.h).

- The function GETFINGERPRINT takes takes two **long**s, *iterations* and *startSeed* (*i.e.*, as a proxy for the packet body), and returns a **long**, which is the checksum, or fingerprint, of the packet body.

You should not need to modify this code for these assignments. If you believe that you do, please start a discussion on piazza, or see the instructors to discuss why you think it should be done.

Your code should be configurable by the following input parameters (`argc, argv`):

- n: the number of threads. Note that one thread is the dispatcher, so there will be $n-1$ sources and $n-1$ worker threads due to the fixed mapping of work to worker—(NUMSOURCES in the code).

- T: the total number of packets from each source—(NUMPACKETS in the code).

- D: the number of entries in each Lamport queue—(QUEUEDEPTH in the code).

- W: the expected amount of work per packet—(MEAN in the code). median, not mean

Next, you will perform the following set of experiments across various cross-products of these parameters. Each data point is a measurement taken on a dynamic system (a computer...) and is thus subject to noise. As a result, some care should be taken to extract representative data—we would propose running some reasonable number of experiments for each data point and selecting the median value as the representative. For Uniform packets, something like 5 data points should suffice whereas for Exponentially Distributed packets 11 may be required to get relatively smooth plots—please use your own engineering judgment to decide how many trials you require. You should set the *seed* parameter for the PACKETSOURCE object to the trial number (or some deterministic function thereof) to ensure that you're seeing a reasonable variation in load from each source. In each of the following experiments, we describe a plot that you should produce, analyze and discuss in the writeup.

1. **Parallel Overhead** Run SERIAL and SERIAL-QUEUE ($D = 32$) on uniformly distributed packets with $W \in \{25, 50, 100, 200, 400, 800\}$ and $n \in \{2, 9, 14\}$, corresponding to 1, 8, and 13 sources. Let $T \approx \frac{2^{24}}{nW}$ for each experiment. Each should take about a second. Plot the speedup (*i.e.* ratio of SERIAL-QUEUE runtime to SERIAL runtime) for each n (*i.e.* one curve for each n) on the Y-axis vs. W on the X-axis. Using the SERIAL-QUEUE data, derive the **Worker Rate**, the packet rate of a worker, given W.

2. **Dispatcher Rate** Run PARALLEL ($D = 32$) on uniformly distributed packets with $W = 1$, $n \in \{2, 3, 5, 9, 14, 28\}$, $T = \frac{2^{20}}{n-1}$. Plot the ratio of $(n-1)T$ to runtime of PARALLEL (*i.e.* packets per second) on the Y-axis vs. n on the X-axis.

3. **Speedup with Constant Load** Run SERIAL and PARALLEL ($D = 32$) on constnat load packets with $W \in \{1000, 2000, 4000, 8000\}$ and $n \in \{2, 3, 5, 9, 14, 28\}$. Plot the speedup of PARALLEL to SERIAL for each W (*i.e.* a curve for each load) on the Y-axis vs. the number of cores, n, on the X-axis. How does the expected speedup compare with the measured speedup?

4. **Speedup with Uniform Load** Run SERIAL and PARALLEL ($D = 32$) on uniformly distributed packets with $W \in \{1000, 2000, 4000, 8000\}$ and $n \in \{2, 3, 5, 9, 14, 28\}$. Let $T = 2^{17}$. Plot the speedup of PARALLEL to SERIAL for each W (*i.e.* a curve for each expected load) on the Y-axis vs. the number of cores, n, on the X-axis. How does the expected speedup compare with the measured speedup?

5. **Speedup with Exponentially Distributed Load** Run the same experiment as **Speedup with uniform load**, except use the exponentially distributed packets from the packet generator. How do these results compare with those with the uniform packets?

Design and Test Document

The design for this problem is more complicated than the previous assignment. You need to spend some time thinking of how you will test your queues, workers, and dispatcher as independent (or semi-independent) entities. You also need to think about how they will work together.

Also, the hypotheses are more complicated for this assignment. You should come up with one hypothesis per experiment. You should leave yourself enough time to adjust if your hypothesis does not match the experimental results. This may require beginning the implementation before the design

review. For example, the assignment specifies that you should use a Lamport Queue. Consider starting your queue testing and implememtation early.

The design and test document is due 4/23/2019. Note that this due date (and the due date for the final assignment) supercedes the dates in the syllabus posted on Piazza.

Writeup

Please submit a typeset report summarizing your results from the experiments and the conclusions you draw from them. Your report should include the five plots (graphs) as specified above. Please make these readable, and use your best judgement—if you think it is easier to read by splitting it up into more graphs, that is okay. Also, submit the working code for SERIAL, SERIAL-QUEUE and PARALLEL, as well as the test code. Both the writeup and code should be submitted through svn.